
Enterprise Data Ops Airflow
Presented by Nam Tonthat
2023-06-21

📔 Agenda

- 🚀 CI/CD
  - Present
  - Planned
- 🌈 Future of Airflow
- ❓ Questions

Present CI/CD

- Builds a new image for any changes
- Publishes this image `~500 MB` each time on updated changes
- Takes 10 minutes for each build

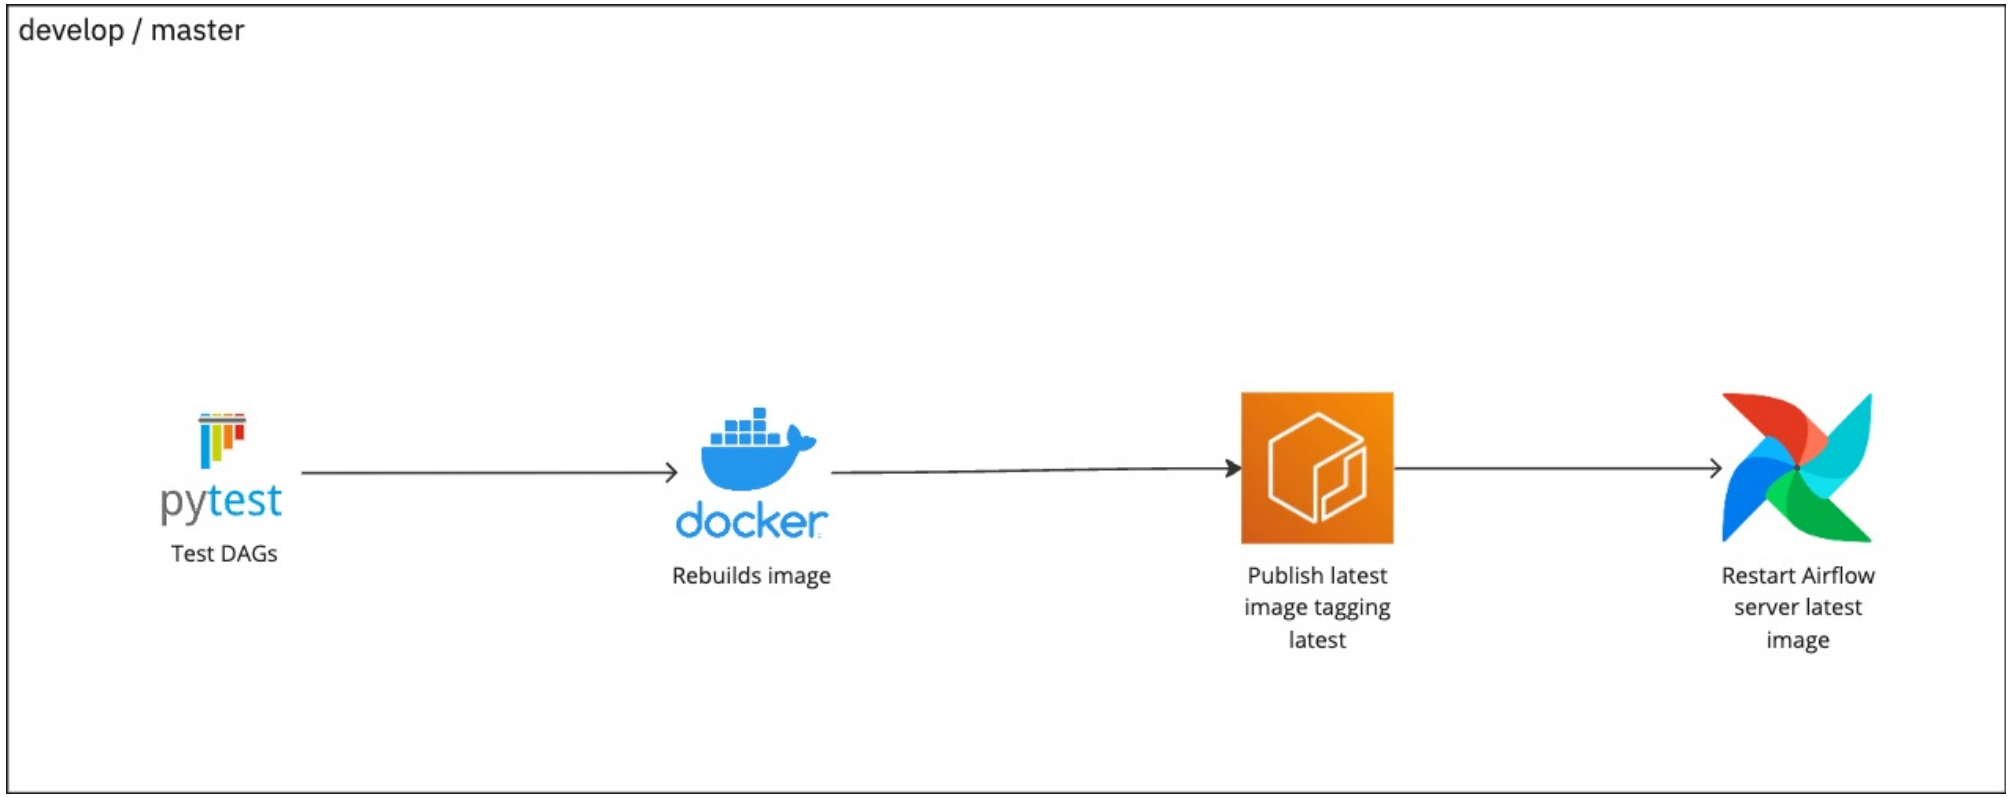

🎯 Planned CI/CD

- Goals:
  - Reduce the time to build by 50%
- Achieved by:
  i. Sync `DAGs` / `variables` into `ga-data-airflow-dags-{env}` S3 bucket.
  ii. Splitting the Dockerfile into a multi stage build
    - `airflow_base` : for requirements
    - `airflow_release` : for `entrypoint.sh` / `airflow.cfg` changes
  iii. Update `entrypoint.sh` using `aws s3 sync` to sync DAGs back
- *Other changes*:
  - Using SEEK's open source software plugins (`seek-oss/docker-ecr-publish`)

🎯 Planned CI/CD - `develop`

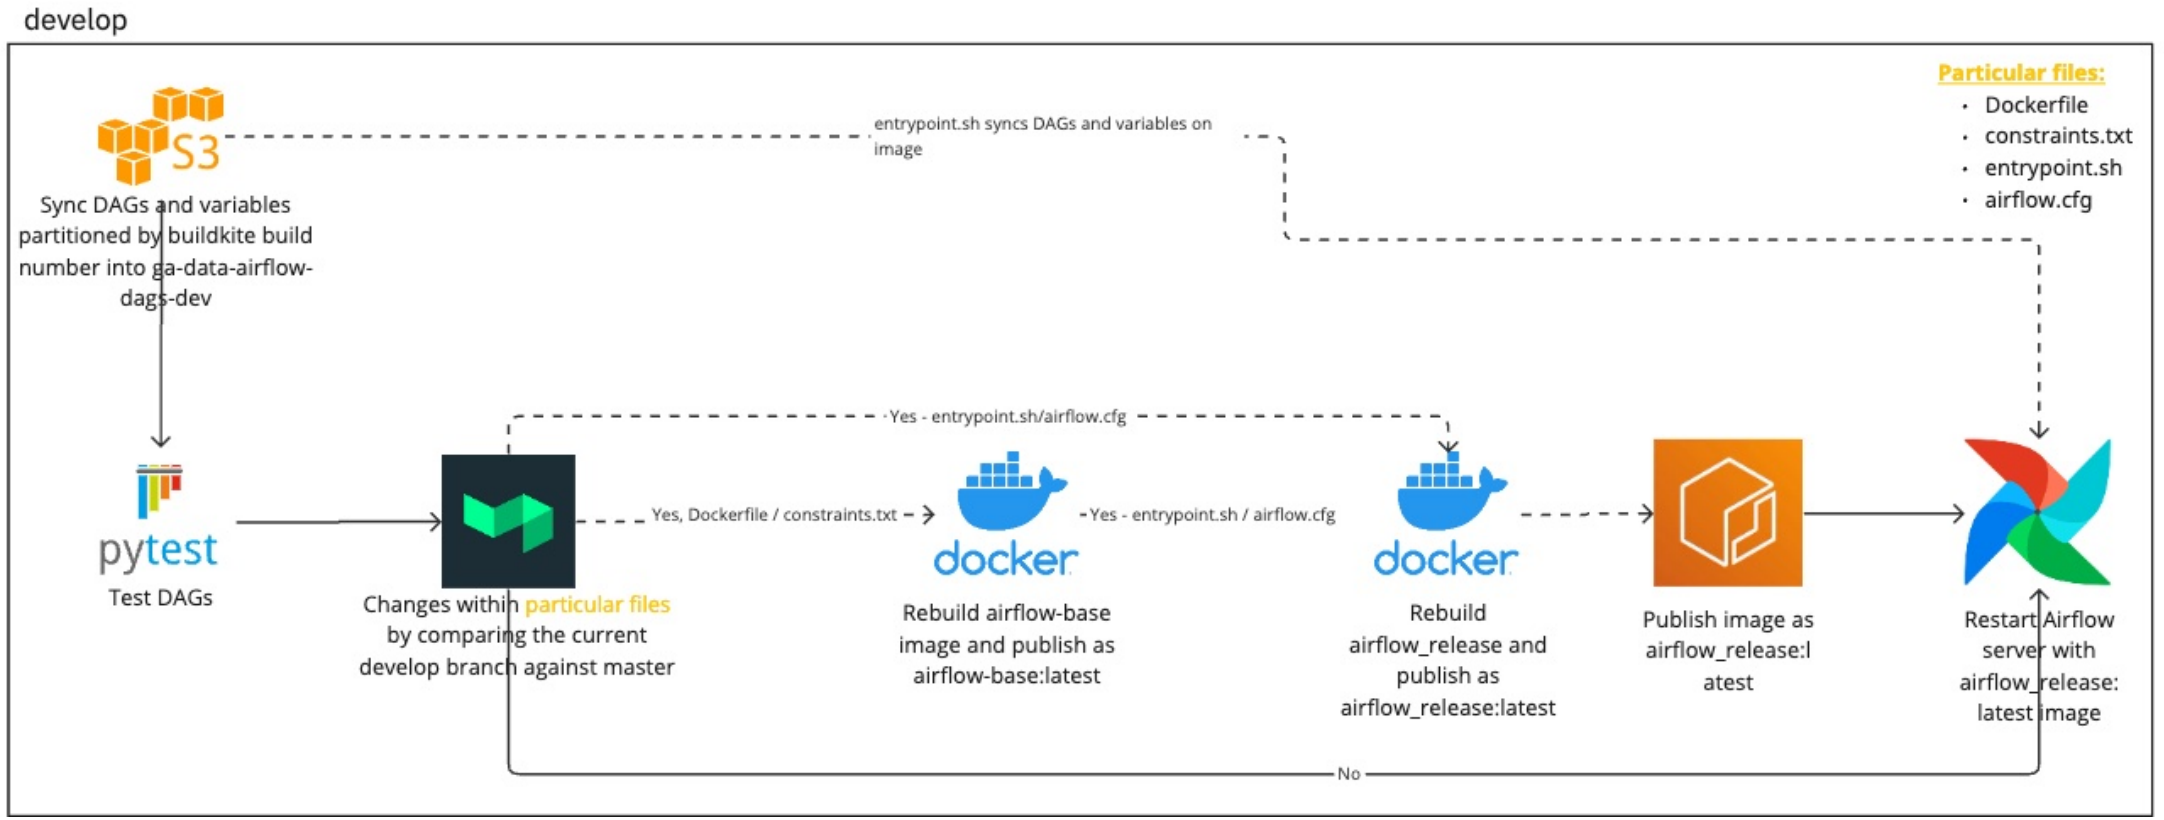

🎯 Planned CI/CD - `master`

Difference between `develop` and `master` is that each PR is a **squash and merge.**

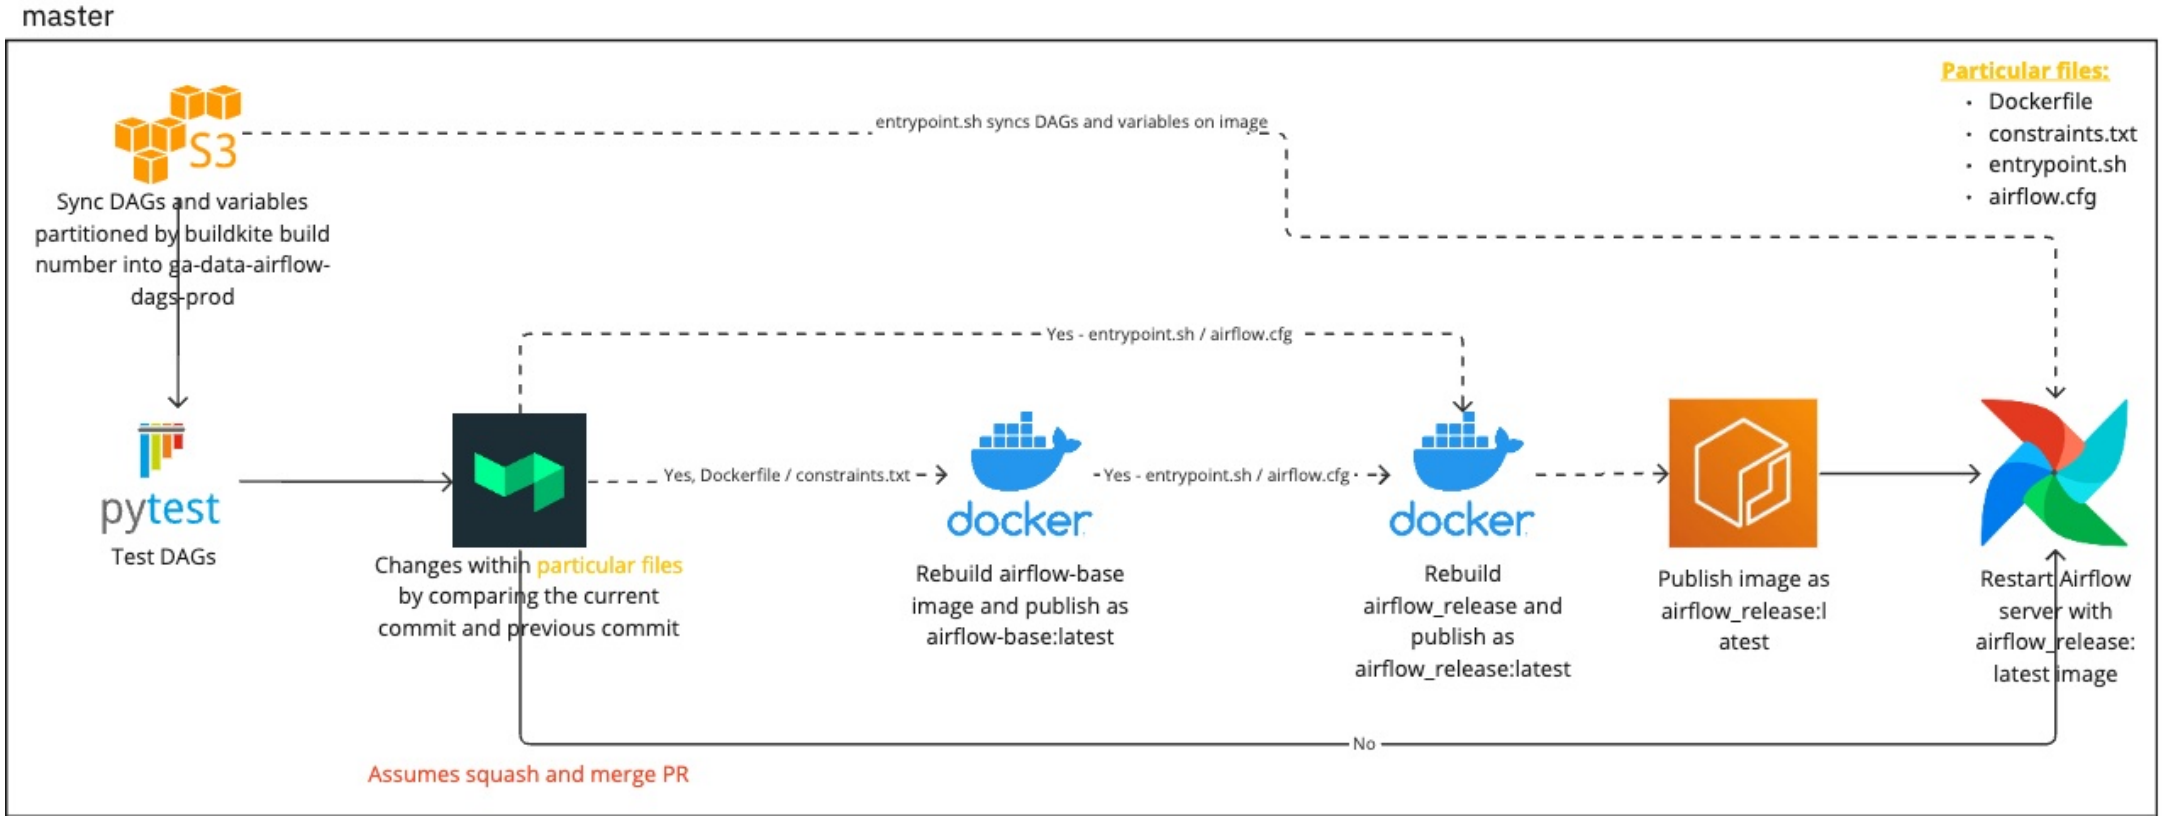

Timing differences - Present CI/CD

- Build time: 10 minutes

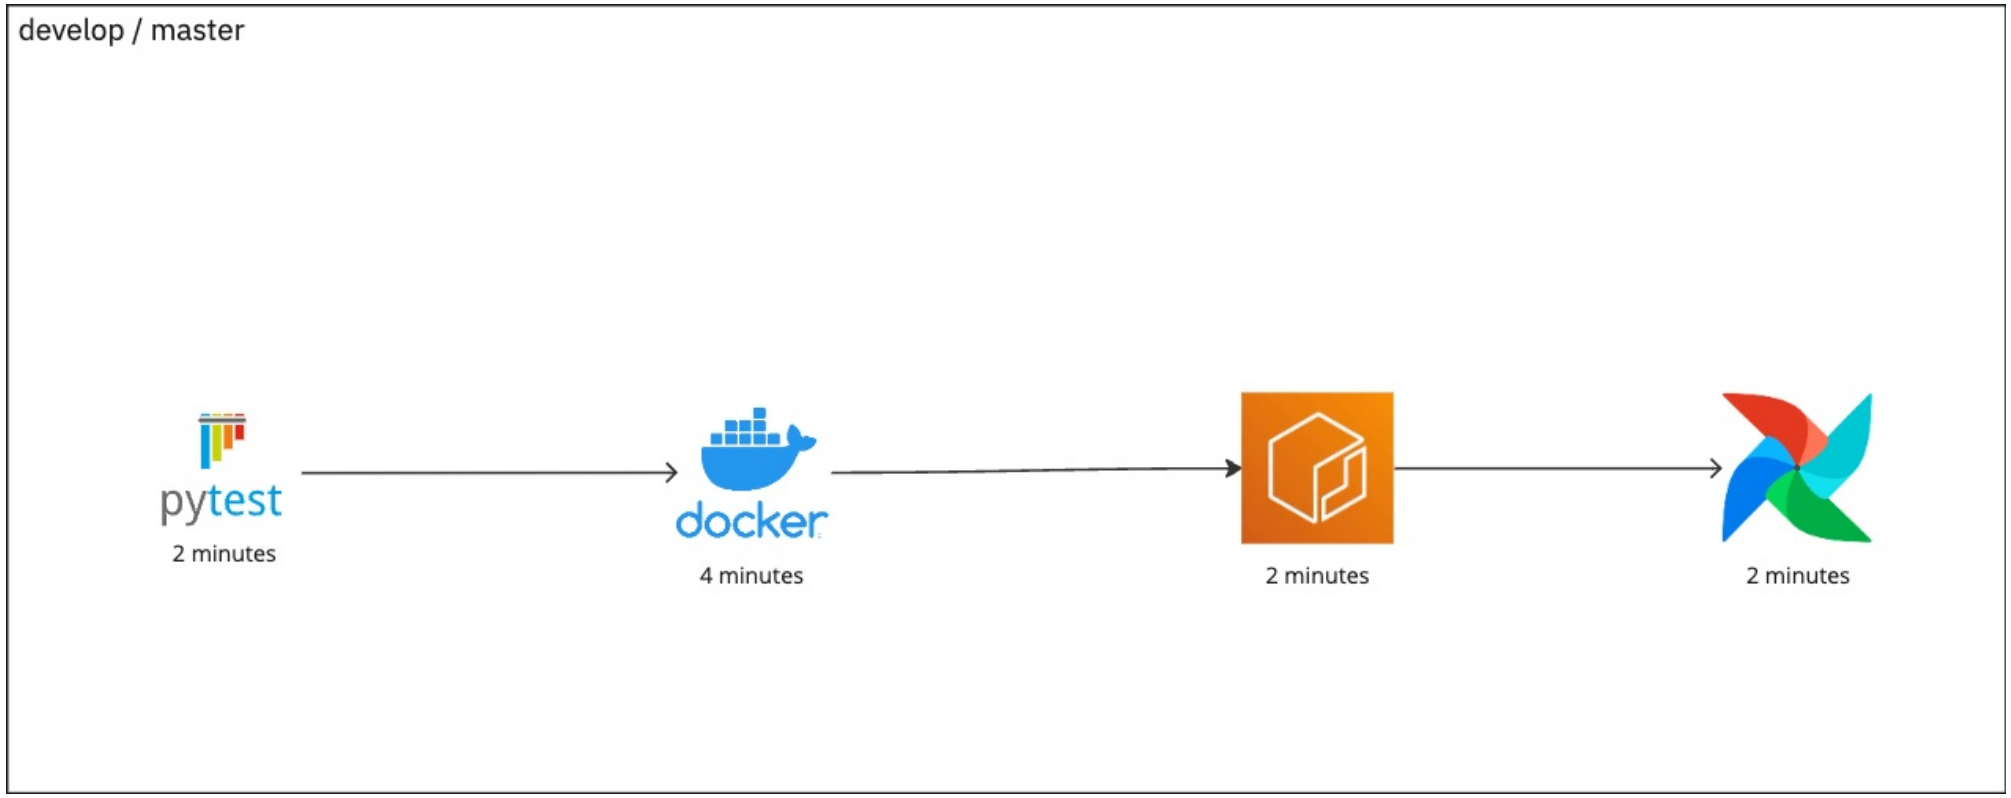

Timing differences - Planned CI/CD

- Build times:
  - 5 minutes for DAG / variable changes (most of the time ~95%)
  - 10 minutes for rebuilding image

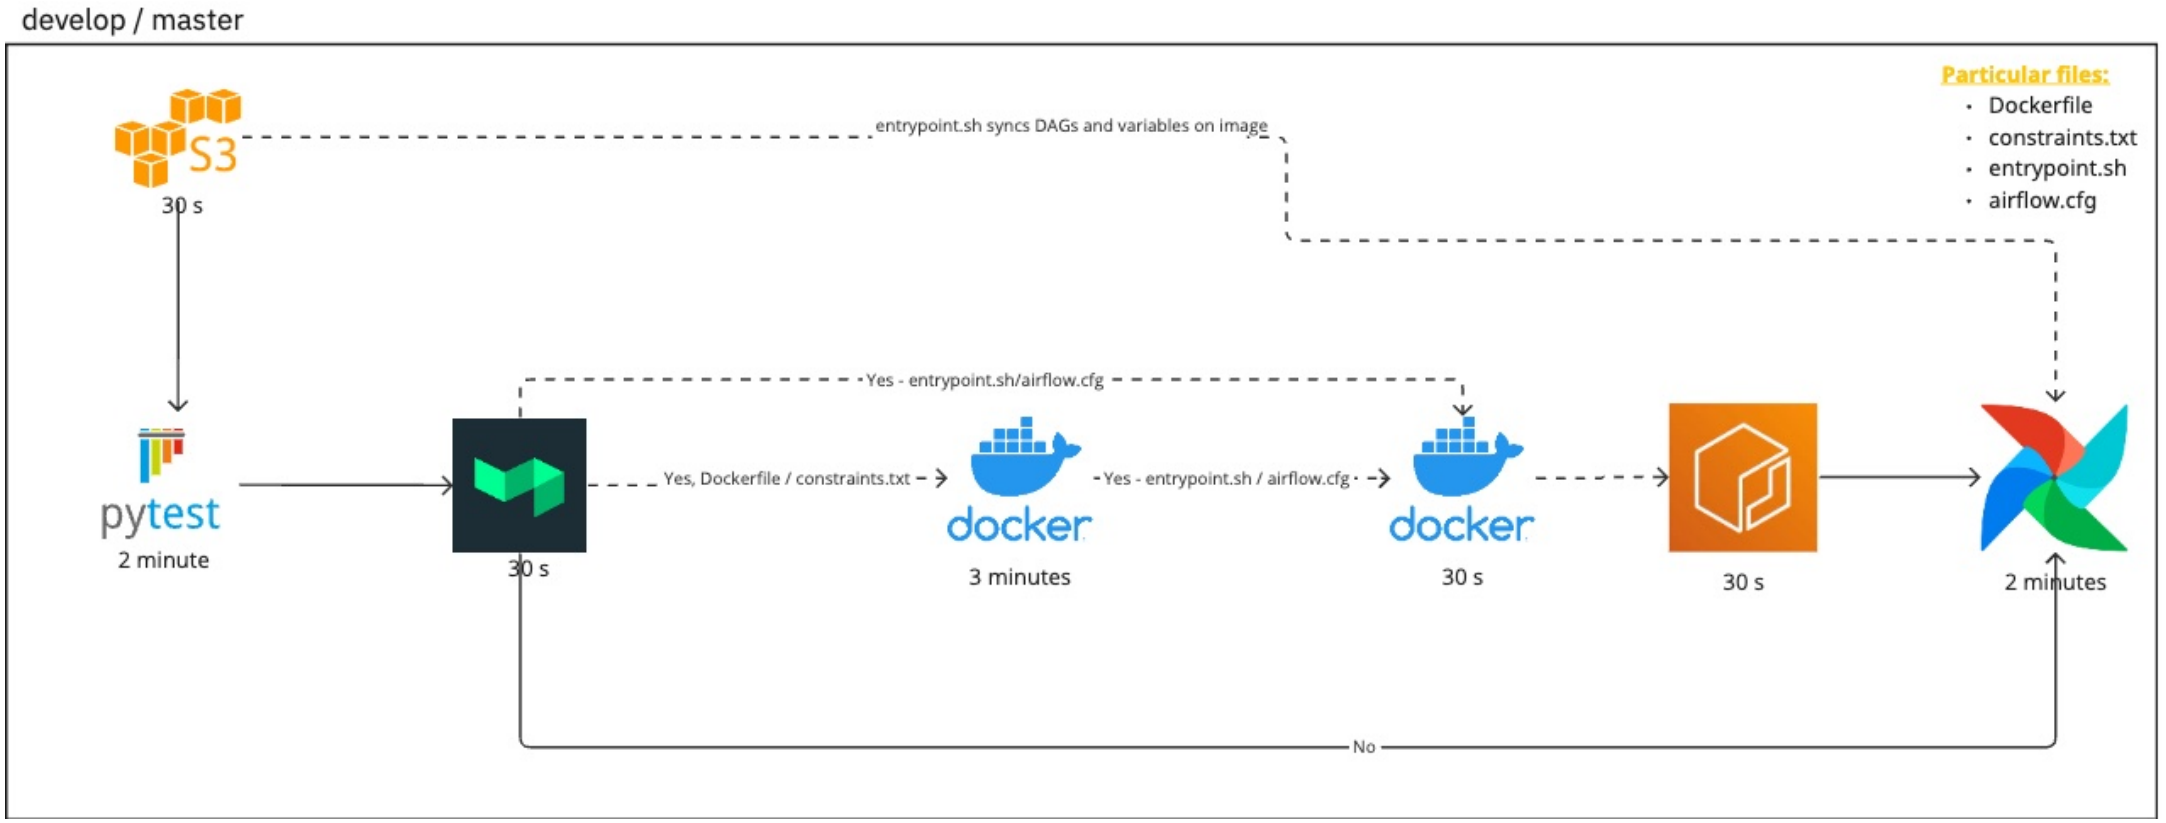

? Questions